Counting Triangles & The Curse of the Last Reducer

Siddharth Suri

Sergei Vassilvitskii

Yahoo! Research

Why Count Triangles?

Sergei Vassilvitskii

Why Count Triangles?

Clustering Coefficient:

Given an undirected graph $G = (V, E)$

cc(v) = fraction of v's neighbors who are neighbors themselves

$$= \frac{|\{(u, w) \in E | u \in \Gamma(v) \wedge w \in \Gamma(v)\}|}{\binom{d_v}{2}}$$

Clustering Coefficient:

Given an undirected graph $G = (V, E)$

cc(v) = fraction of v's neighbors who are neighbors themselves

$$= \frac{|\{(u, w) \in E | u \in \Gamma(v) \wedge w \in \Gamma(v)\}|}{\binom{d_v}{2}}$$

cc (●) = N/A

cc (●) = 1/3

cc (●) = 1

cc (●) = 1

Clustering Coefficient:

Given an undirected graph $G = (V, E)$

cc(v) = fraction of v's neighbors who are neighbors themselves

$$= \frac{|\{(u,w) \in E | u \in \Gamma(v) \wedge w \in \Gamma(v)\}|}{\binom{d_v}{2}} = \frac{\#\Delta s \text{ incident on } v}{\binom{d_v}{2}}$$

cc (🔵) = N/A

cc (🔴) = 1/3

cc (🟢) = 1

cc (⚫) = 1

Captures how tight-knit the network is around a node.

cc (🔴) = 0.5

cc (🔵) = 0.1

vs.

6

Captures how tight-knit the network is around a node.

cc (🔴) = 0.5   vs.   cc (🔵) = 0.1

Network Cohesion:

    - Tightly knit communities foster more trust, social norms. [Coleman '88, Portes '88]

Structural Holes:

    - Individuals benefit form bridging [Burt '04, '07]

Why MapReduce?

De facto standard for parallel computation on large data

- Widely used at: Yahoo!, Google, Facebook,

- Also at: New York Times, Amazon.com, Match.com, …

- Commodity hardware

- Reliable infrastructure

- Data continues to outpace available RAM !

Sergei Vassilvitskii

Sequential Version:

```
foreach v in V
    foreach u,w in Adjacency(v)
        if (u,w) in E
            Triangles[v]++
```

v

Triangles[v]=0

Sequential Version:

```
foreach v in V
    foreach u,w in Adjacency(v)
        if (u,w) in E
            Triangles[v]++
```


Triangles[v]=1

How to Count Triangles

Sequential Version:

```
foreach v in V
    foreach u,w in Adjacency(v)
        if (u,w) in E
            Triangles[v]++
```

Triangles[v]=1

How to Count Triangles

Sequential Version:

```
foreach v in V
     foreach u,w in Adjacency(v)
         if (u,w) in E
             Triangles[v]++
```

Running time: $\sum\limits_{v \in V} d_v^2$

Even for sparse graphs can be quadratic if one vertex has high degree.

Parallelize the edge checking phase

Parallelize the edge checking phase

– Map 1: For each v send $(v, \Gamma(v))$ to single machine.

– Reduce 1: Input: $\langle v; \Gamma(v) \rangle$

Output: all 2 paths $\langle (v_1, v_2); u \rangle$ where $v_1, v_2 \in \Gamma(u)$

(●,●);●   (●,●); ●   (● ,●);●

Parallelize the edge checking phase

- Map 1: For each v send $(v, \Gamma(v))$ to single machine.
- Reduce 1: Input: $\langle v; \Gamma(v) \rangle$

  Output: all 2 paths $\langle (v_1, v_2); u \rangle$ where $v_1, v_2 \in \Gamma(u)$

  (●,●);●   (●,●); ●   (● ,●);●

- Map 2: Send $\langle (v_1, v_2); u \rangle$ and $\langle (v_1, v_2); \$ \rangle$ for $(v_1, v_2) \in E$ to same machine.

- Reduce 2: input: $\langle (v, w); u_1, u_2, \ldots, u_k, \$? \rangle$

  Output: if $\$$ part of the input, then: $u_i = u_i + {}^1\!/\!_3$

  (●,●);●, $\$$ $\longrightarrow$ ● $+{}^1\!/\!_3$ ● $+{}^1\!/\!_3$ ● $+{}^1\!/\!_3$

  (●,●);● $\longrightarrow$

15

Data skew

How much parallelization can we achieve?

- Generate all the paths to check in parallel
- The running time becomes $\max\limits_{v \in V} d_v^2$

Sergei Vassilvitskii

Data skew

How much parallelization can we achieve?

- Generate all the paths to check in parallel

- The running time becomes $\max\limits_{v \in V} d_v^2$

Naive parallelization does not help with data skew

– Some nodes will have very high degree

– Example. 3.2 Million followers, must generate 10 Trillion (10^13) potential edges to check.

– Even if generating 100M edges per second, 100K seconds ~ 27 hours.

"Just 5 more minutes"

Running the naive algorithm on LiveJournal Graph

– 80% of reducers done after 5 min

– 99% done after 35 min

Adapting the Algorithm

Approach 1: Dealing with skew directly

- currently every triangle counted 3 times (once per vertex)

- Running time quadratic in the degree of the vertex

- Idea: Count each once, from the perspective of lowest degree vertex

- Does this heuristic work?

Adapting the Algorithm

Approach 1: Dealing with skew directly

– currently every triangle counted 3 times (once per vertex)

– Running time quadratic in the degree of the vertex

– Idea: Count each once, from the perspective of lowest degree vertex

– Does this heuristic work?

Approach 2: Divide & Conquer

– Equally divide the graph between machines

– But any edge partition will be bound to miss triangles

– Divide into overlapping subgraphs, account for the overlap

How to Count Triangles Better

Sequential Version [Schank '07]:

```
foreach v in V
    foreach u,w in Adjacency(v)
        if deg(u) > deg(v) && deg(w) > deg(v)
            if (u,w) in E
                Triangles[v]++
```

Sergei Vassilvitskii

Does it make a difference?

Why does it help?

- Partition nodes into two groups:
  - Low: $\mathcal{L} = \{v : d_v \leq \sqrt{m}\}$
  - High: $\mathcal{H} = \{v : d_v > \sqrt{m}\}$
- There are at most n low nodes; each produces at most $O(m)$ paths
- There are at most $2\sqrt{m}$ high nodes
  - Each produces paths to other high nodes: $O(m)$ paths per node

Dealing with Skew

Why does it help?

- Partition nodes into two groups:
  - Low: $\mathcal{L} = \{v : d_v \leq \sqrt{m}\}$
  - High: $\mathcal{H} = \{v : d_v > \sqrt{m}\}$
- There are at most n low nodes; each produces at most $O(m)$ paths
- There are at most $2\sqrt{m}$ high nodes
  - Each produces paths to other high nodes: $O(m)$ paths per node

- These two are identical !
- Therefore, no mapper can produce substantially more work than others.
- Total work is $O(m^{3/2})$, which is optimal

Approach 2: Graph Split

Partitioning the nodes:

- Previous algorithm shows one way to achieve better parallelization
- But what if even $O(m)$ is too much. Is it possible to divide input into smaller chunks?

Graph Split Algorithm:

- Partition vertices into p equal sized groups $V_1, V_2, \ldots, V_p$.
- Consider all possible triples (V_i, V_j, V_k) and the induced subgraph:
$$G_{ijk} = G\left[V_i \cup V_j \cup V_k\right]$$
- Compute the triangles on each G_{ijk} separately.

Some Triangles present in multiple subgraphs:

in p-2 subgraphs

in 1 subgraph

in ~p² subgraphs

Can count exactly how many subgraphs each triangle will be in

Approach 2: Graph Split

Analysis:

- Each subgraph has $O(m/p^2)$ edges in expectation.
- Very balanced running times

Distribution of Reducer Completion Times

27

Approach 2: Graph Split

Analysis:

- Very balanced running times
- p controls memory needed per machine

Reducer Memory vs. Amount Shuffled (LiveJournal)

Approach 2: Graph Split

Analysis:

– Very balanced running times

– p controls memory needed per machine

– Total work: $p^3 \cdot O((m/p^2)^{3/2}) = O(m^{3/2})$, independent of p

Approach 2: Graph Split

Analysis:

- Very balanced running times
- p controls memory needed per machine
- Total work: $p^3 \cdot O((m/p^2)^{3/2}) = O(m^{3/2})$, independent of p

Input too big: paging

Shuffle time increases with duplication

Runtime vs. ρ for LiveJournal

Overall

Naive Parallelization Doesn't help with Data Skew

Sergei Vassilvitskii

Related Work

- ## Tsourakakis et al. [09]:

  - Count global number of triangles by estimating the trace of the cube of the matrix

  - Don't specifically deal with skew, obtain high probability approximations.

- ## Becchetti et al. [08]

  - Approximate the number of triangles per node

  - Use multiple passes to obtain a better and better approximation

Conclusions

Conclusions

Think about data skew.... and avoid the curse

Sergei Vassilvitskii

Conclusions

Think about data skew…. and avoid the curse

- Get programs to run faster

Conclusions

Think about data skew.... and avoid the curse

- Get programs to run faster
- Publish more papers

33

Conclusions

Think about data skew.... and avoid the curse

- – Get programs to run faster

- – Publish more papers

- – Get more sleep

33

Conclusions

Think about data skew…. and avoid the curse

- Get programs to run faster

- Publish more papers

- Get more sleep

- ..

Sergei Vassilvitskii

Conclusions

Think about data skew…. and avoid the curse

- Get programs to run faster

- Publish more papers

- Get more sleep

- ..

- The possibilities are endless!

Thank You